

```
In [1]: #Practical No 1
        #ajit waman B54
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        %matplotlib inline
        #so that we can view the graphs inside the notebook
```

```
In [2]: #from google.colab import drive
        #drive.mount('/content/gdrive')
```

```
In [3]: s1 = pd.Series(range(1,10,1))
```

```
In [4]: s1
```

```
Out[4]: 0    1
        1    2
        2    3
        3    4
        4    5
        5    6
        6    7
        7    8
        8    9
        dtype: int64
```

```
In [5]: s3 = pd.Series({1:21, 2:13,3:45})
```

```
In [6]: s3
```

```
Out[6]: 1    21
        2    13
        3    45
        dtype: int64
```

```
In [7]: s2 = pd.Series([1, 2, 3, 4], index=['p', 'q', 'r', 's'], name='one')
```

```
In [8]: s2
```

```
Out[8]: p    1
        q    2
        r    3
        s    4
        Name: one, dtype: int64
```

```
In [9]: df1 = pd.DataFrame(s2)
df1
```

Out[9]:

	one
p	1
q	2
r	3
s	4

```
In [11]: df2 = pd.read_csv("/home/kj-comp/california_housing_test.csv")
#dataframe_name = pd.read_<format>(filename)
```

```
In [12]: df2.head(10)
```

Out[12]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households
0	-122.05	37.37	27.0	3885.0	661.0	1537.0	606.0
1	-118.30	34.26	43.0	1510.0	310.0	809.0	277.0
2	-117.81	33.78	27.0	3589.0	507.0	1484.0	495.0
3	-118.36	33.82	28.0	67.0	15.0	49.0	11.0
4	-119.67	36.33	19.0	1241.0	244.0	850.0	237.0
5	-119.56	36.51	37.0	1018.0	213.0	663.0	204.0
6	-121.43	38.63	43.0	1009.0	225.0	604.0	218.0
7	-120.65	35.48	19.0	2310.0	471.0	1341.0	441.0
8	-122.84	38.40	15.0	3080.0	617.0	1446.0	599.0
9	-118.02	34.08	31.0	2402.0	632.0	2830.0	603.0

```
In [13]: df2.tail(3)
```

Out[13]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househo
2997	-119.70	36.30	10.0	956.0	201.0	693.0	22
2998	-117.12	34.10	40.0	96.0	14.0	46.0	1
2999	-119.63	34.42	42.0	1765.0	263.0	753.0	26

```
In [14]: df2['median_house_value_new']=df2['median_house_value']+111
```

In [15]: `df2.tail(3)`

Out[15]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househo
2997	-119.70	36.30	10.0	956.0	201.0	693.0	22
2998	-117.12	34.10	40.0	96.0	14.0	46.0	1
2999	-119.63	34.42	42.0	1765.0	263.0	753.0	26

In [16]: `# write`
`# <dataframe's name>.to_<file_format>(<file_name>)`

In [17]: `df2.to_json('data1.json')`

In [18]: `len(df2['total_rooms'])`

Out[18]: 3000

In [19]: `df2['total_rooms'].count()`

Out[19]: 3000

In [20]: `df2['total_rooms'].mean()`

Out[20]: 2599.5786666666667

In [21]: `df2['total_rooms'].sum()`

Out[21]: 7798736.0

In [22]: `df2['total_rooms'].median()`

Out[22]: 2106.0

In [23]: `df2['total_rooms'].std()`

Out[23]: 2155.59333162558

In [24]: `df2['total_rooms'].min()`

Out[24]: 6.0

In [25]: `df2['total_rooms'].max()`

Out[25]: 30450.0

```
In [26]: df2['total_rooms'].describe()
```

```
Out[26]: count      3000.000000
mean       2599.578667
std        2155.593332
min         6.000000
25%        1401.000000
50%        2106.000000
75%        3129.000000
max        30450.000000
Name: total_rooms, dtype: float64
```

```
In [27]: df2['total_rooms'].cumsum()
```

```
Out[27]: 0          3885.0
1          5395.0
2          8984.0
3          9051.0
4         10292.0
...
2995      7790662.0
2996      7795919.0
2997      7796875.0
2998      7796971.0
2999      7798736.0
Name: total_rooms, Length: 3000, dtype: float64
```

```
In [28]: # When you give the whole dataframe, then all numerical columns will be analysed
df2.mean()
```

```
Out[28]: longitude      -119.589200
latitude       35.635390
housing_median_age    28.845333
total_rooms       2599.578667
total_bedrooms     529.950667
population       1402.798667
households        489.912000
median_income       3.807272
median_house_value  205846.275000
median_house_value_new 205957.275000
dtype: float64
```

In [29]: `df2.describe()`

Out[29]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
count	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000
mean	-119.589200	35.63539	28.845333	2599.578667	529.950667	1402.798667
std	1.994936	2.12967	12.555396	2155.593332	415.654368	1030.543012
min	-124.180000	32.56000	1.000000	6.000000	2.000000	5.000000
25%	-121.810000	33.93000	18.000000	1401.000000	291.000000	780.000000
50%	-118.485000	34.27000	29.000000	2106.000000	437.000000	1155.000000
75%	-118.020000	37.69000	37.000000	3129.000000	636.000000	1742.750000
max	-114.490000	41.92000	52.000000	30450.000000	5419.000000	11935.000000

In [31]: `df = pd.read_csv("/home/kj-comp/california_housing_test.csv")`

In [32]: `df.describe()`

Out[32]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
count	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000	3000.000000
mean	-119.589200	35.63539	28.845333	2599.578667	529.950667	1402.798667
std	1.994936	2.12967	12.555396	2155.593332	415.654368	1030.543012
min	-124.180000	32.56000	1.000000	6.000000	2.000000	5.000000
25%	-121.810000	33.93000	18.000000	1401.000000	291.000000	780.000000
50%	-118.485000	34.27000	29.000000	2106.000000	437.000000	1155.000000
75%	-118.020000	37.69000	37.000000	3129.000000	636.000000	1742.750000
max	-114.490000	41.92000	52.000000	30450.000000	5419.000000	11935.000000

In [33]: `df.columns`

Out[33]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income', 'median_house_value'], dtype='object')

In [34]: `df['longitude']`

Out[34]:

0	-122.05
1	-118.30
2	-117.81
3	-118.36
4	-119.67
...	...
2995	-119.86
2996	-118.14
2997	-119.70
2998	-117.12
2999	-119.63

Name: longitude, Length: 3000, dtype: float64

In [35]: `df.longitude`

Out[35]:

0	-122.05
1	-118.30
2	-117.81
3	-118.36
4	-119.67
...	...
2995	-119.86
2996	-118.14
2997	-119.70
2998	-117.12
2999	-119.63

Name: longitude, Length: 3000, dtype: float64

In [36]: `df.iloc[:,1:3]`

Out[36]:

	latitude	housing_median_age
0	37.37	27.0
1	34.26	43.0
2	33.78	27.0
3	33.82	28.0
4	36.33	19.0
...
2995	34.42	23.0
2996	34.06	27.0
2997	36.30	10.0
2998	34.10	40.0
2999	34.42	42.0

3000 rows × 2 columns

```
In [38]: # importing pandas as pd
import pandas as pd
# making data frame from csv file
data = pd.read_csv("employees.csv")
data.head(10)
```

Out[38]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	Douglas	Male	8/6/1993	12:42 PM	97308.0	6.945	True	Marketing
1	Thomas	Male	3/31/1996	6:53 AM	61933.0	4.170	True	NaN
2	Maria	Female	4/23/1993	11:17 AM	130590.0	11.858	False	Finance
3	Jerry	Male	3/4/2005	1:00 PM	138705.0	9.340	True	Finance
4	Larry	Male	1/24/1998	4:47 PM	101004.0	1.389	True	Client Services
5	Dennis	Male	4/18/1987	1:35 AM	115163.0	10.125	False	Legal
6	Ruby	Female	8/17/1987	4:20 PM	65476.0	10.012	True	Product
7	NaN	Female	7/20/2015	10:43 AM	45906.0	11.598	NaN	Finance
8	Angela	Female	11/22/2005	6:29 AM	95570.0	18.523	True	Engineering
9	Frances	Female	8/8/2002	6:51 AM	139852.0	7.524	True	Business Development

```
In [39]: data.describe()
```

Out[39]:

	Salary	Bonus %
count	969.000000	1000.000000
mean	90579.972136	10.207555
std	32916.214577	5.528481
min	35013.000000	1.015000
25%	62666.000000	5.401750
50%	90370.000000	9.838500
75%	118733.000000	14.838000
max	149908.000000	19.944000

In [40]: data.isnull()

Out[40]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	True
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
995	False	True	False	False	False	False	False	False
996	False	False	False	False	False	False	False	False
997	False	False	False	False	False	False	False	False
998	False	False	False	False	False	False	False	False
999	False	False	False	False	False	False	False	False

1000 rows × 8 columns

In [41]: data.notnull()

Out[41]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	False
2	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True
...
995	True	False	True	True	True	True	True	True
996	True	True	True	True	True	True	True	True
997	True	True	True	True	True	True	True	True
998	True	True	True	True	True	True	True	True
999	True	True	True	True	True	True	True	True

1000 rows × 8 columns


```
In [42]: data.isnull().sum()
```

```
Out[42]: First Name      67  
Gender      145  
Start Date    0  
Last Login Time  0  
Salary       31  
Bonus %       0  
Senior Management  67  
Team         43  
dtype: int64
```

```
In [43]: # filling a null values using fillna()  
data["Gender"].fillna("No Gender", inplace = True)
```

```
In [44]: data.isnull().sum()
```

```
Out[44]: First Name      67  
Gender      0  
Start Date    0  
Last Login Time  0  
Salary       31  
Bonus %       0  
Senior Management  67  
Team         43  
dtype: int64
```

```
In [45]: # will replace Nan value in dataframe with value -99
import numpy as np
data.replace(to_replace = np.nan, value = -99)
```

Out[45]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	Douglas	Male	8/6/1993	12:42 PM	97308.0	6.945	True	Marketing
1	Thomas	Male	3/31/1996	6:53 AM	61933.0	4.170	True	-99
2	Maria	Female	4/23/1993	11:17 AM	130590.0	11.858	False	Finance
3	Jerry	Male	3/4/2005	1:00 PM	138705.0	9.340	True	Finance
4	Larry	Male	1/24/1998	4:47 PM	101004.0	1.389	True	Client Services
...
995	Henry	No Gender	11/23/2014	6:09 AM	132483.0	16.655	False	Distribution
996	Phillip	Male	1/31/1984	6:30 AM	42392.0	19.675	False	Finance
997	Russell	Male	5/20/2013	12:39 PM	96914.0	1.421	False	Product
998	Larry	Male	4/20/2013	4:45 PM	60500.0	11.985	False	Business Development
999	Albert	Male	5/15/2012	6:24 PM	129949.0	10.169	True	Sales

1000 rows × 8 columns

In [46]: *# filling a missing value with previous ones*
 data.fillna(method = 'pad')

Out[46]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	Douglas	Male	8/6/1993	12:42 PM	97308.0	6.945	True	Marketing
1	Thomas	Male	3/31/1996	6:53 AM	61933.0	4.170	True	Marketing
2	Maria	Female	4/23/1993	11:17 AM	130590.0	11.858	False	Finance
3	Jerry	Male	3/4/2005	1:00 PM	138705.0	9.340	True	Finance
4	Larry	Male	1/24/1998	4:47 PM	101004.0	1.389	True	Client Services
...
995	Henry	No Gender	11/23/2014	6:09 AM	132483.0	16.655	False	Distribution
996	Phillip	Male	1/31/1984	6:30 AM	42392.0	19.675	False	Finance
997	Russell	Male	5/20/2013	12:39 PM	96914.0	1.421	False	Product
998	Larry	Male	4/20/2013	4:45 PM	60500.0	11.985	False	Business Development
999	Albert	Male	5/15/2012	6:24 PM	129949.0	10.169	True	Sales

1000 rows × 8 columns

```
In [47]: # filling a missing value with previous ones
data.fillna(method = 'pad')
```

Out[47]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	Senior Management	Team
0	Douglas	Male	8/6/1993	12:42 PM	97308.0	6.945	True	Marketing
1	Thomas	Male	3/31/1996	6:53 AM	61933.0	4.170	True	Marketing
2	Maria	Female	4/23/1993	11:17 AM	130590.0	11.858	False	Finance
3	Jerry	Male	3/4/2005	1:00 PM	138705.0	9.340	True	Finance
4	Larry	Male	1/24/1998	4:47 PM	101004.0	1.389	True	Client Services
...
995	Henry	No Gender	11/23/2014	6:09 AM	132483.0	16.655	False	Distribution
996	Phillip	Male	1/31/1984	6:30 AM	42392.0	19.675	False	Finance
997	Russell	Male	5/20/2013	12:39 PM	96914.0	1.421	False	Product
998	Larry	Male	4/20/2013	4:45 PM	60500.0	11.985	False	Business Development
999	Albert	Male	5/15/2012	6:24 PM	129949.0	10.169	True	Sales

1000 rows × 8 columns

```
In [48]: data.dropna(axis=1)
```

Out[48]:

	Gender	Start Date	Last Login Time	Bonus %
0	Male	8/6/1993	12:42 PM	6.945
1	Male	3/31/1996	6:53 AM	4.170
2	Female	4/23/1993	11:17 AM	11.858
3	Male	3/4/2005	1:00 PM	9.340
4	Male	1/24/1998	4:47 PM	1.389
...
995	No Gender	11/23/2014	6:09 AM	16.655
996	Male	1/31/1984	6:30 AM	19.675
997	Male	5/20/2013	12:39 PM	1.421
998	Male	4/20/2013	4:45 PM	11.985
999	Male	5/15/2012	6:24 PM	10.169

1000 rows × 4 columns

```
In [49]: # importing pandas as pd
import pandas as pd
# Creating the dataframe
df = pd.DataFrame({"A": [12, 4, 5, None, 1],
                  "B": [None, 2, 54, 3, None],
                  "C": [20, 16, None, 3, 8],
                  "D": [14, 3, None, None, 6]})
# Print the dataframe
df
```

Out[49]:

	A	B	C	D
0	12.0	NaN	20.0	14.0
1	4.0	2.0	16.0	3.0
2	5.0	54.0	NaN	NaN
3	NaN	3.0	3.0	NaN
4	1.0	NaN	8.0	6.0

```
In [50]: df.interpolate(method='linear', limit_direction='forward')
```

Out[50]:

	A	B	C	D
0	12.0	NaN	20.0	14.0
1	4.0	2.0	16.0	3.0
2	5.0	54.0	9.5	4.0
3	3.0	3.0	3.0	5.0
4	1.0	3.0	8.0	6.0

```
In [51]: #remove white space everywhere
text="today is Monday"
#df['Col Name'] = df['Col Name'].str.replace(' ', '')
text.replace(' ', '')
```

Out[51]: 'todayisMonday'

```
In [52]: text=' Today'
text.lstrip()
```

Out[52]: 'Today'

```
In [53]: text='Today '
text.rstrip()
```

Out[53]: 'Today'

```
In [54]: text=' Today '
text.strip()
```

Out[54]: 'Today'

```
In [55]: import pandas
import scipy
import numpy
from sklearn.preprocessing import MinMaxScaler
# data values
X = [ [110, 200], [120, 800], [310, 400], [140, 900], [510, 200], [653, 400] ,
      [310, 880]
    ]
# transform data
scaler = MinMaxScaler(feature_range=(0,5))
rescaledX = scaler.fit_transform(X)
```

```
In [56]: X
```

```
Out[56]: [[110, 200],
          [120, 800],
          [310, 400],
          [140, 900],
          [510, 200],
          [653, 400],
          [310, 880]]
```

```
In [57]: rescaledX
```

```
Out[57]: array([[0.          , 0.          ],
                [0.09208103, 4.28571429],
                [1.84162063, 1.42857143],
                [0.27624309, 5.          ],
                [3.68324125, 0.          ],
                [5.          , 1.42857143],
                [1.84162063, 4.85714286]])
```

```
In [58]: from sklearn.preprocessing import StandardScaler
import pandas
import numpy
# data values
X = [ [110, 200], [120, 800], [310, 400], [140, 900], [510, 200], [653, 400] ,
      [310, 880]
    ]
# scaler
scaler = StandardScaler().fit(X)
rescaledX = scaler.transform(X)
```

```
In [59]: rescaledX
```

```
Out[59]: array([[ -1.02004521, -1.17792918],
                [-0.96841602,  0.90076937],
                [ 0.01253852, -0.48502966],
                [-0.86515765,  1.24721913],
                [ 1.04512224, -1.17792918],
                [ 1.78341961, -0.48502966],
                [ 0.01253852,  1.17792918]])
```

```
In [60]: from sklearn.preprocessing import Normalizer
import pandas
import numpy
# data values
X = [ [110, 200], [120, 800], [310, 400], [140, 900], [510, 200], [653, 400] ,
      [310, 880]
      ] ]
# normalize values
scaler = Normalizer().fit(X)
normalizedX = scaler.transform(X)
```

```
In [61]: normalizedX
```

```
Out[61]: array([[0.48191875, 0.87621591],
                [0.14834045, 0.98893635],
                [0.61257167, 0.79041505],
                [0.15370701, 0.98811647],
                [0.9309732 , 0.36508753],
                [0.8527326 , 0.52234769],
                [0.33225942, 0.94318804]])
```

```
In [62]: from sklearn.preprocessing import Binarizer
import pandas
import numpy
# data values
X = [ [501, 200], [120, 800], [310, 400], [140, 900], [510, 200], [653, 400] ,
      [310, 880]
      ] ]
# binarize data
binarizer = Binarizer(threshold=500).fit(X)
binaryX = binarizer.transform(X)
```

```
In [63]: binaryX
```


```
Out[63]: array([[1, 0],
                [0, 1],
                [0, 0],
                [0, 1],
                [1, 0],
                [1, 0],
                [0, 1]])
```

```
In [72]: import pandas as pd
import numpy as np
# Read in the CSV file and convert "?" to NaN
headers = ["symboling", "normalized_losses", "make", "fuel_type", "aspiration",
"num_doors", "body_style", "drive_wheels", "engine_location",
"wheel_base", "length", "width", "height", "curb_weight",
"engine_type", "num_cylinders", "engine_size", "fuel_system",
"bore", "stroke", "compression_ratio", "horsepower", "peak_rpm",
"city_mpg", "highway_mpg", "price"]
df = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/au
tos/imports-85.data", header=None, names=headers, na_values="?")
df.head()
```

Out[72]:

	symboling	normalized_losses	make	fuel_type	aspiration	num_doors	body_style	drive_wheels
0	3	NaN	alfa-romero	gas	std	two	convertible	
1	3	NaN	alfa-romero	gas	std	two	convertible	
2	1	NaN	alfa-romero	gas	std	two	hatchback	
3	2	164.0	audi	gas	std	four	sedan	
4	2	164.0	audi	gas	std	four	sedan	

5 rows × 26 columns



In [73]: df.dtypes

```
Out[73]: symboling          int64
normalized_losses    float64
make                 object
fuel_type            object
aspiration            object
num_doors             object
body_style            object
drive_wheels          object
engine_location        object
wheel_base            float64
length                float64
width                 float64
height                float64
curb_weight           int64
engine_type            object
num_cylinders          object
engine_size           int64
fuel_system            object
bore                  float64
stroke                float64
compression_ratio      float64
horsepower             float64
peak_rpm              float64
city_mpg              int64
highway_mpg           int64
price                 float64
dtype: object
```

In [74]: obj_df = df.select_dtypes(include=['object']).copy()
obj_df.head()

Out[74]:

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_ty
0	alfa-romero	gas	std	two	convertible	rwd	front	dc
1	alfa-romero	gas	std	two	convertible	rwd	front	dc
2	alfa-romero	gas	std	two	hatchback	rwd	front	of
3	audi	gas	std	four	sedan	fwd	front	c
4	audi	gas	std	four	sedan	4wd	front	c

In [75]: `obj_df[obj_df.isnull().any(axis=1)]`

Out[75]:

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_t
27	dodge	gas	turbo	NaN	sedan	fwd	front	
63	mazda	diesel	std	NaN	sedan	fwd	front	

In [76]: `obj_df["num_doors"].value_counts()`

Out[76]:

```
four      114
two        89
Name: num_doors, dtype: int64
```

In [77]: `obj_df = obj_df.fillna({"num_doors": "four"})`

In [78]: `obj_df[obj_df.isnull().any(axis=1)]`

Out[78]:

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_type
--	------	-----------	------------	-----------	------------	--------------	-----------------	-------------

In [79]: `obj_df["num_cylinders"].value_counts()`

Out[79]:

```
four      159
six        24
five       11
eight        5
two          4
twelve       1
three        1
Name: num_cylinders, dtype: int64
```

In [80]: `cleanup_nums = {"num_doors": {"four": 4, "two": 2},
"num_cylinders": {"four": 4, "six": 6, "five": 5, "eight": 8,
"two": 2, "twelve": 12, "three": 3 }}`

```
In [81]: obj_df = obj_df.replace(cleanup_nums)
obj_df.head()
```

```
Out[81]:
```

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_ty
0	alfa-romero	gas	std	2	convertible	rwd	front	dc
1	alfa-romero	gas	std	2	convertible	rwd	front	dc
2	alfa-romero	gas	std	2	hatchback	rwd	front	of
3	audi	gas	std	4	sedan	fwd	front	c
4	audi	gas	std	4	sedan	4wd	front	c

```
In [82]: obj_df.dtypes
```

```
Out[82]: make                object
fuel_type                  object
aspiration                 object
num_doors                  int64
body_style                 object
drive_wheels              object
engine_location            object
engine_type               object
num_cylinders              int64
fuel_system               object
dtype: object
```

```
In [83]: obj_df["body_style"].value_counts()
```

```
Out[83]: sedan                96
hatchback                   70
wagon                       25
hardtop                      8
convertible                   6
Name: body_style, dtype: int64
```

```
In [84]: obj_df["body_style"] = obj_df["body_style"].astype('category')
obj_df.dtypes
```

```
Out[84]: make                object
fuel_type                  object
aspiration                 object
num_doors                  int64
body_style                 category
drive_wheels              object
engine_location            object
engine_type               object
num_cylinders              int64
fuel_system               object
dtype: object
```

```
In [85]: obj_df["body_style_cat"] = obj_df["body_style"].cat.codes
obj_df.head()
```

Out[85]:

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_ty
0	alfa-romero	gas	std	2	convertible	rwd	front	dc
1	alfa-romero	gas	std	2	convertible	rwd	front	dc
2	alfa-romero	gas	std	2	hatchback	rwd	front	of
3	audi	gas	std	4	sedan	fwd	front	c
4	audi	gas	std	4	sedan	4wd	front	c

```
In [86]: pd.get_dummies(obj_df, columns=["drive_wheels"]).head()
```

Out[86]:

	make	fuel_type	aspiration	num_doors	body_style	engine_location	engine_type	num_cylin
0	alfa-romero	gas	std	2	convertible	front	dohc	
1	alfa-romero	gas	std	2	convertible	front	dohc	
2	alfa-romero	gas	std	2	hatchback	front	ohcv	
3	audi	gas	std	4	sedan	front	ohc	
4	audi	gas	std	4	sedan	front	ohc	

```
In [87]: from sklearn.preprocessing import OrdinalEncoder
ord_enc = OrdinalEncoder()
obj_df["make_code"] = ord_enc.fit_transform(obj_df[["make"]])
obj_df[["make", "make_code"]].head(11)
```

Out[87]:

	make	make_code
0	alfa-romero	0.0
1	alfa-romero	0.0
2	alfa-romero	0.0
3	audi	1.0
4	audi	1.0
5	audi	1.0
6	audi	1.0
7	audi	1.0
8	audi	1.0
9	audi	1.0
10	bmw	2.0

```
In [88]: from sklearn.preprocessing import OneHotEncoder
oe_style = OneHotEncoder()
oe_results = oe_style.fit_transform(obj_df[["body_style"]])
pd.DataFrame(oe_results.toarray(), columns=oe_style.categories_).head()
```

Out[88]:

	convertible	hardtop	hatchback	sedan	wagon
0	1.0	0.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0	0.0
2	0.0	0.0	1.0	0.0	0.0
3	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	1.0	0.0

```
In [89]: obj_df = obj_df.join(pd.DataFrame(oe_results.toarray(), columns=oe_style.categories_))
```

```
In [90]: obj_df.head()
```

Out[90]:

	make	fuel_type	aspiration	num_doors	body_style	drive_wheels	engine_location	engine_ty
0	alfa-romero	gas	std	2	convertible	rwd	front	dc
1	alfa-romero	gas	std	2	convertible	rwd	front	dc
2	alfa-romero	gas	std	2	hatchback	rwd	front	ol
3	audi	gas	std	4	sedan	fwd	front	c
4	audi	gas	std	4	sedan	4wd	front	c

```
In [ ]:
```